Parallel DBSCAN

Bowen Chen

DBSCAN Overview

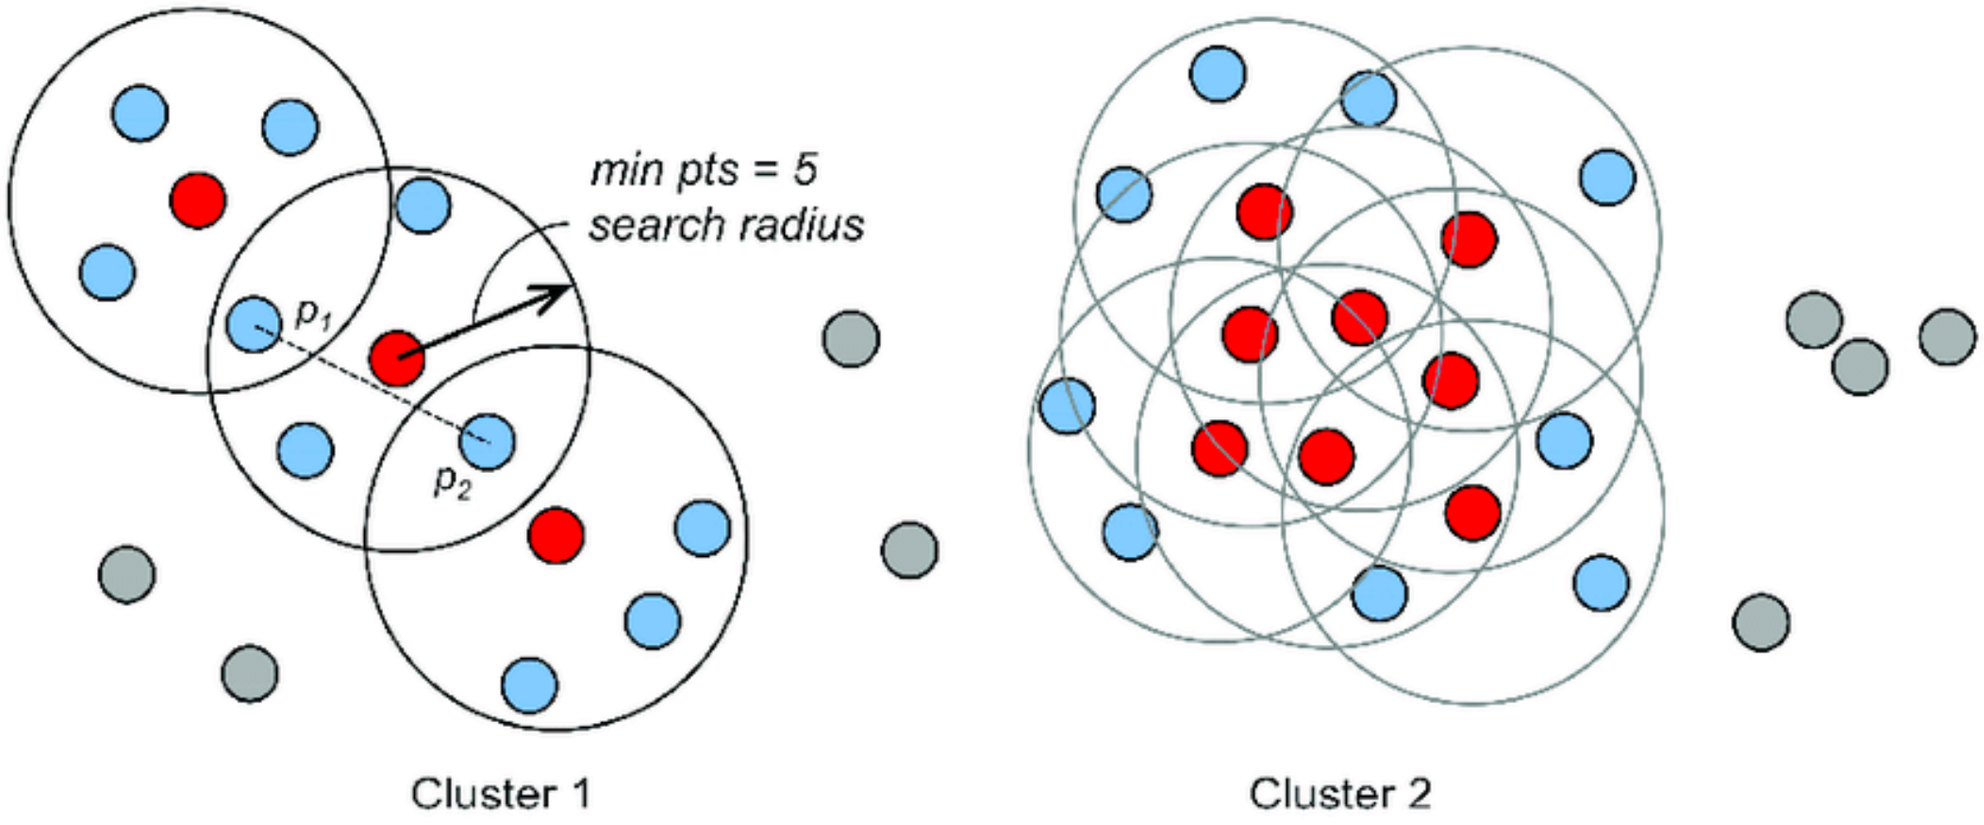

Break down: **1) find core points** **2) growing neighbours**

G-DBSCAN

Graph Construction

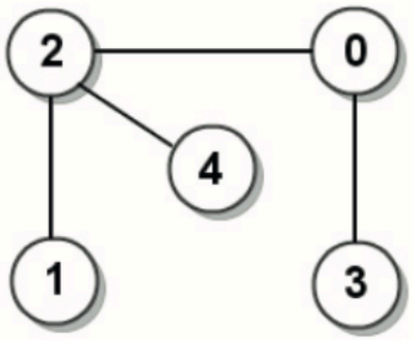

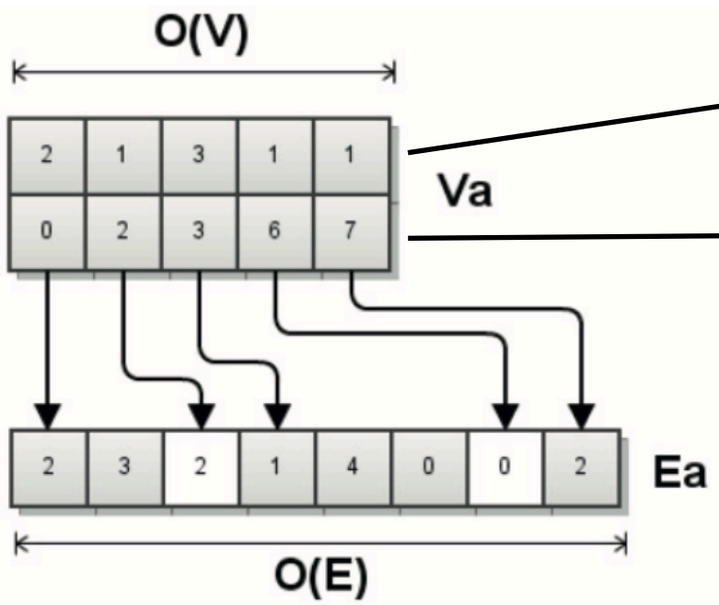

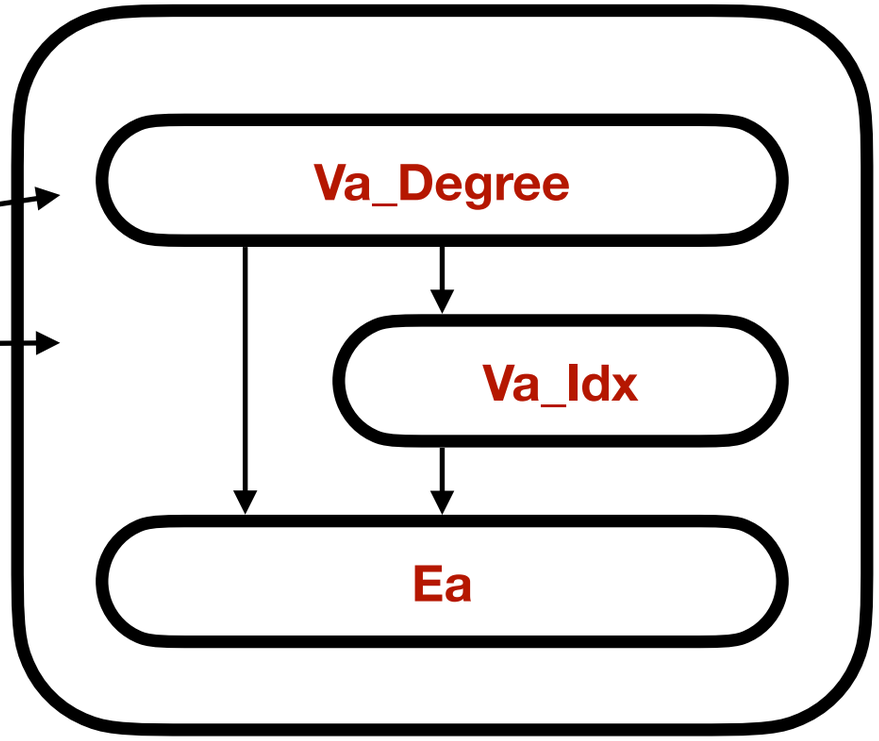

BFS SCAN

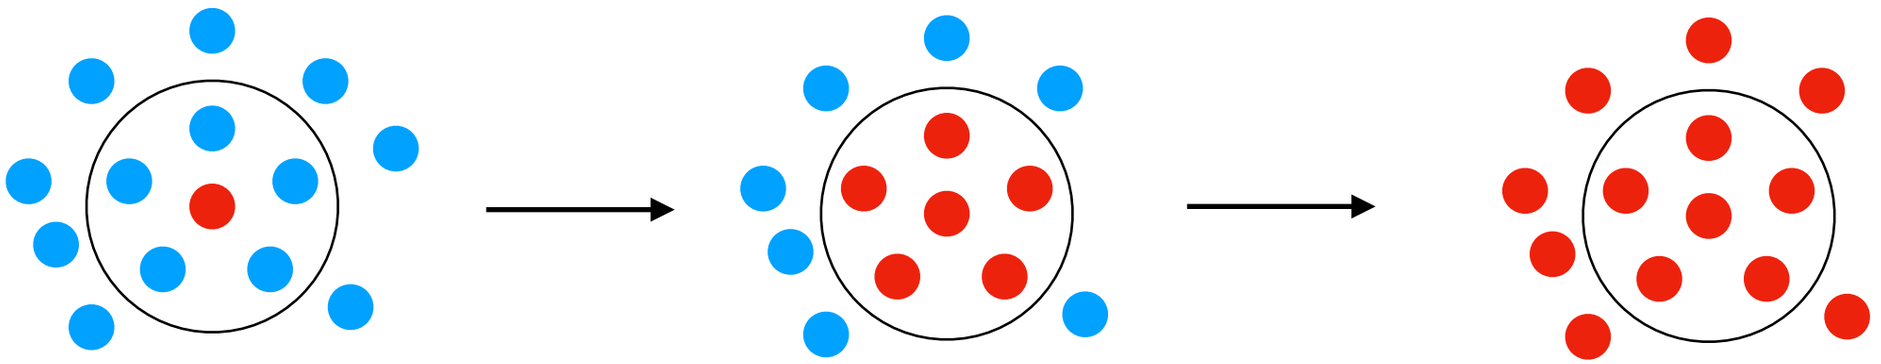

G-DBSCAN

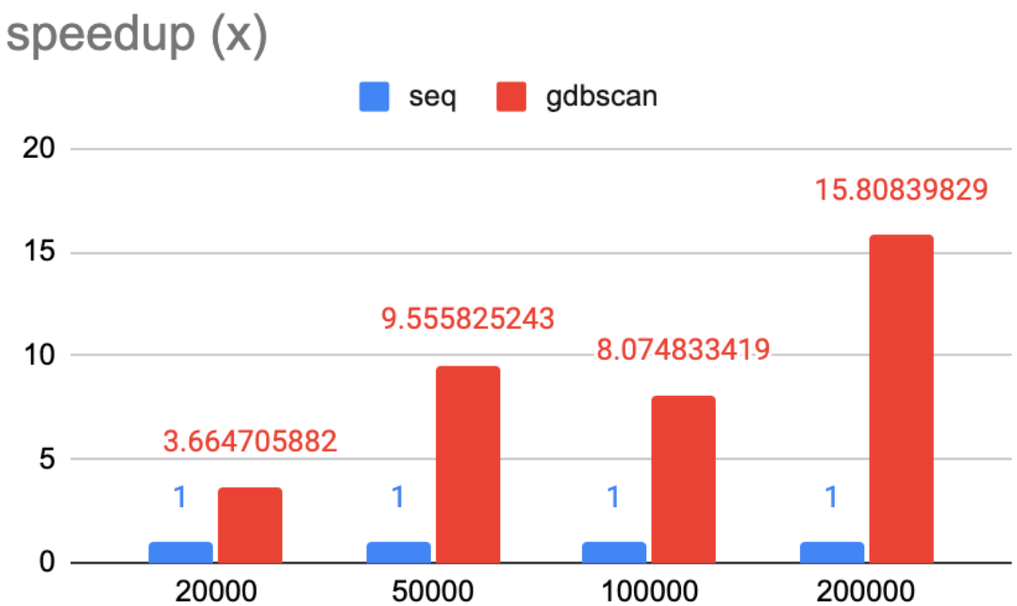

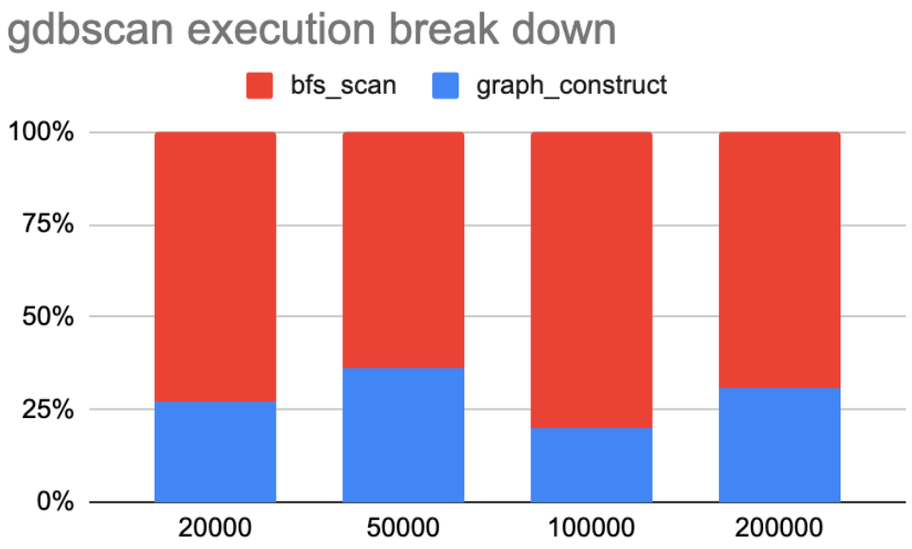

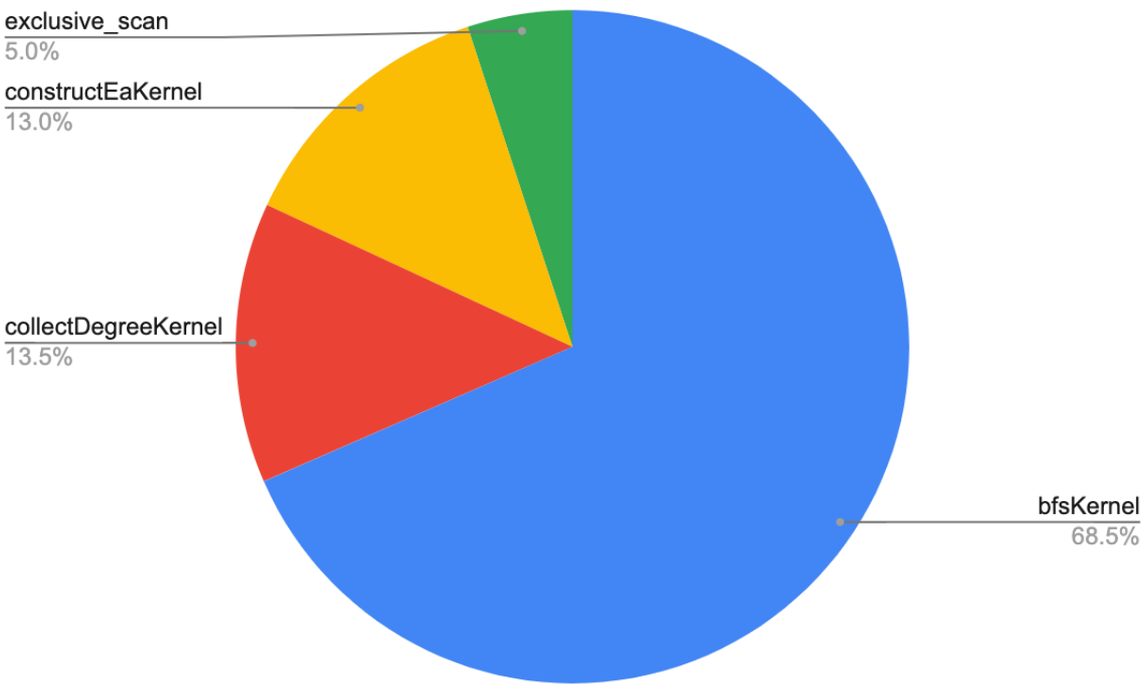

BFS SCAN is the bottleneck!!!

PDSDBSCAN (Disjoint-Set)

Key Idea: postpone union

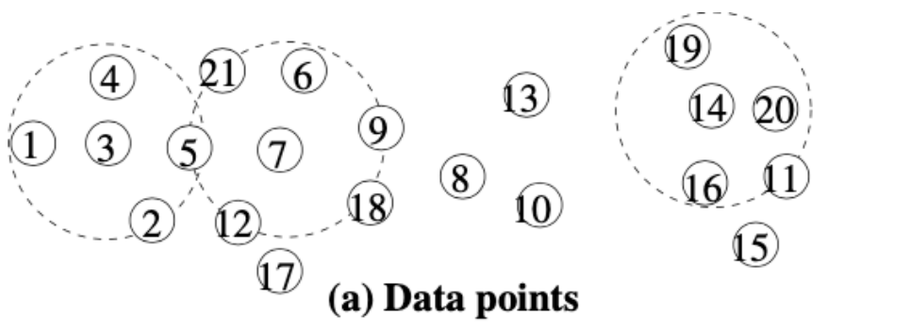

Iterate over points

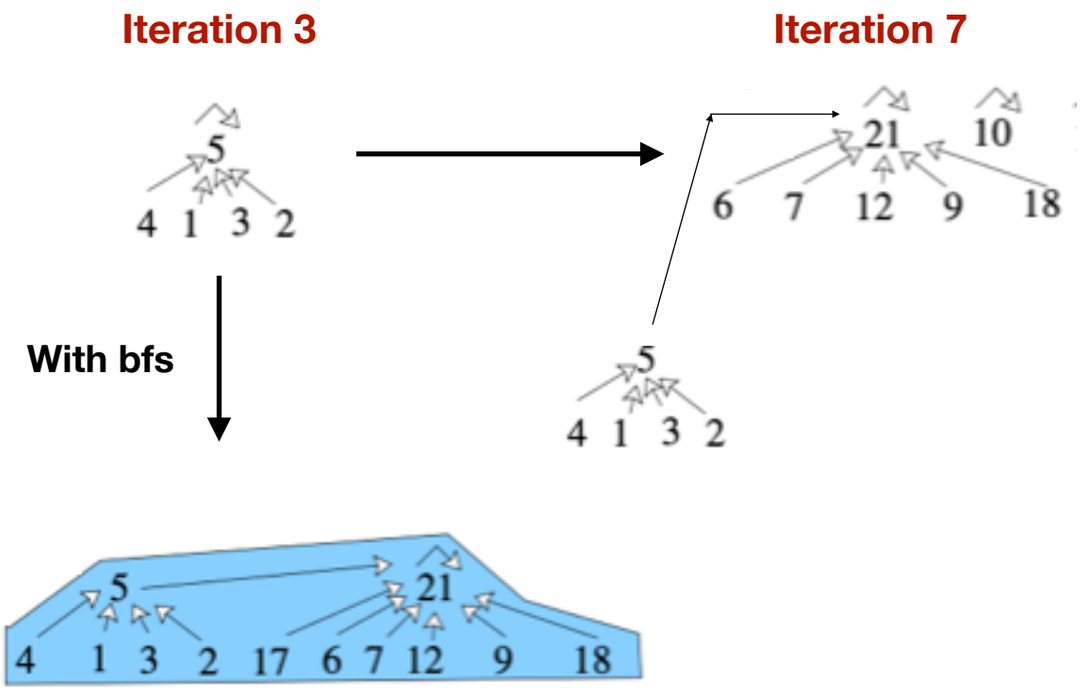

Shared Memory Parallel with omp

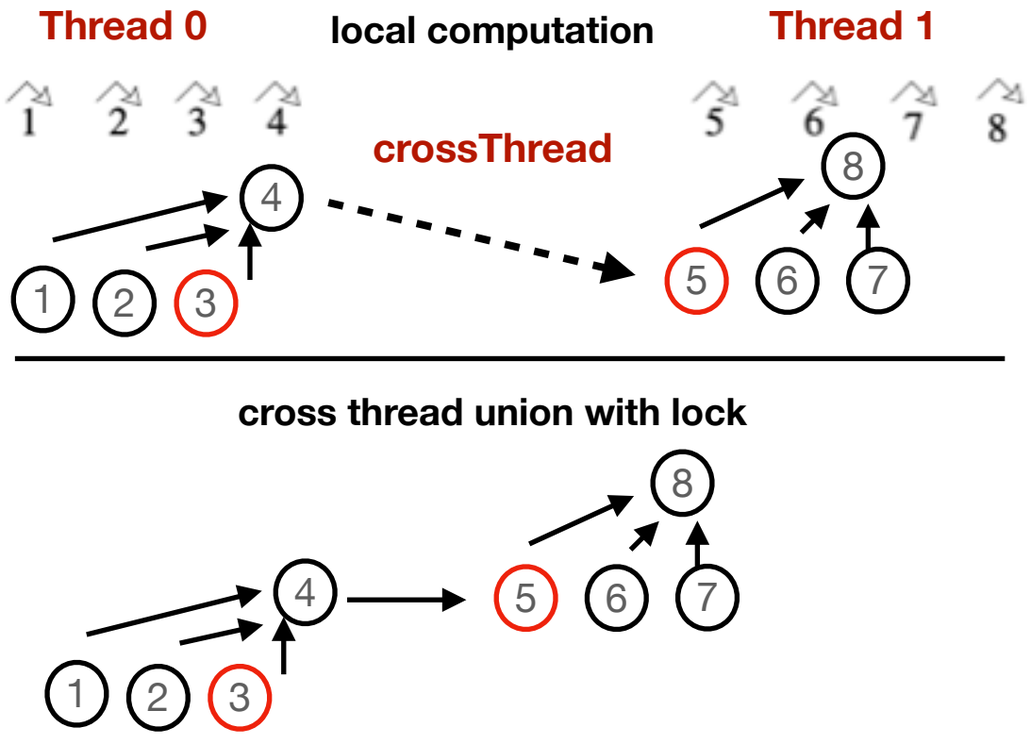

PDSDBSCAN (Disjoint-Set)

Execution time with different threads

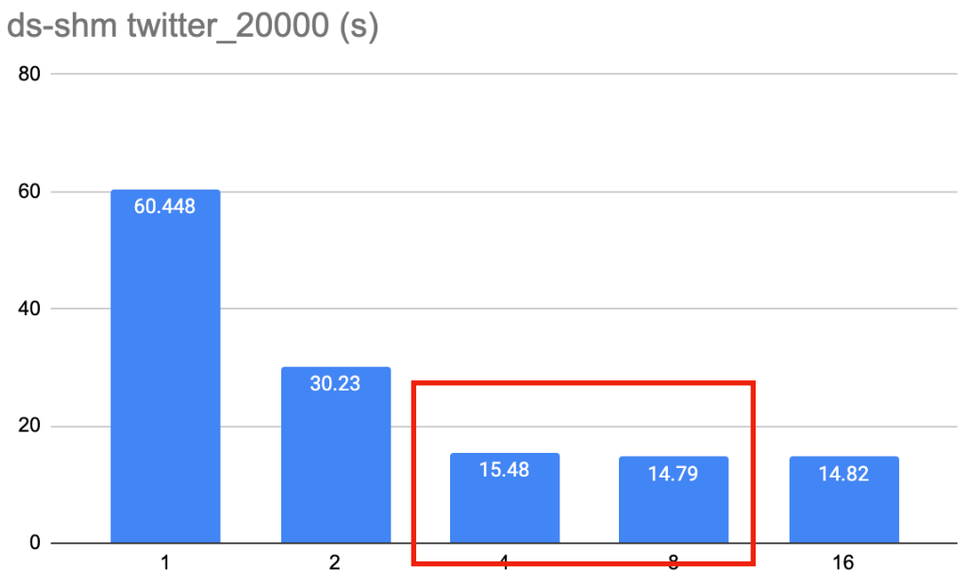

Cache Misses

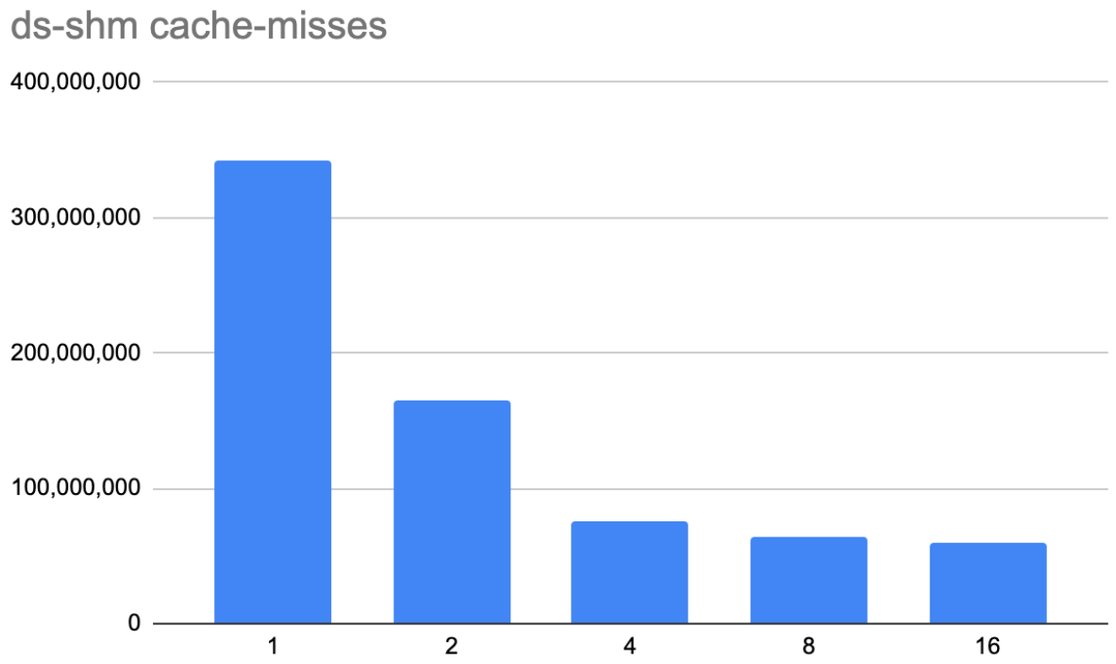

Workload Distribution

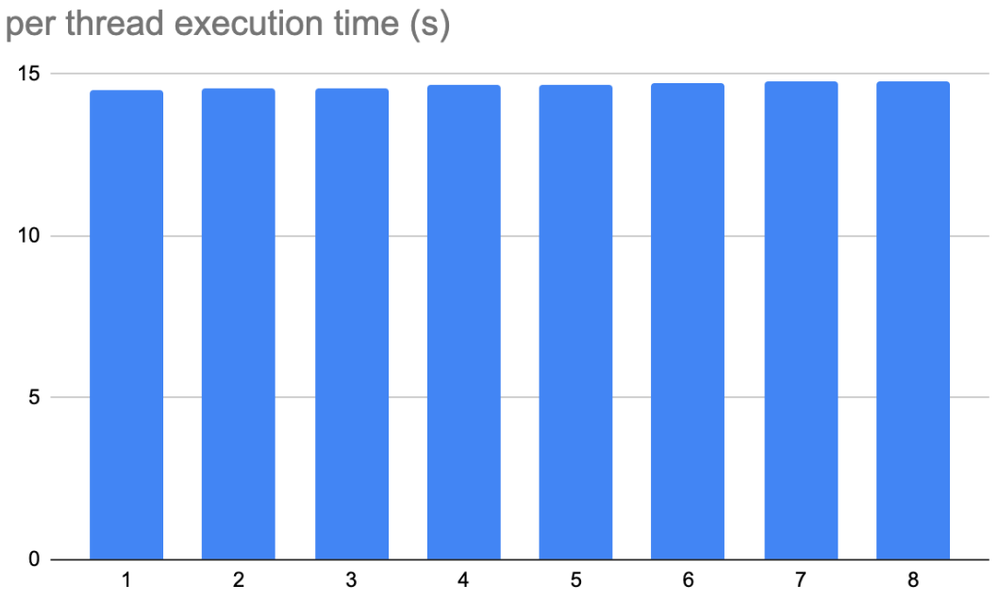

Execution Time Breakdown

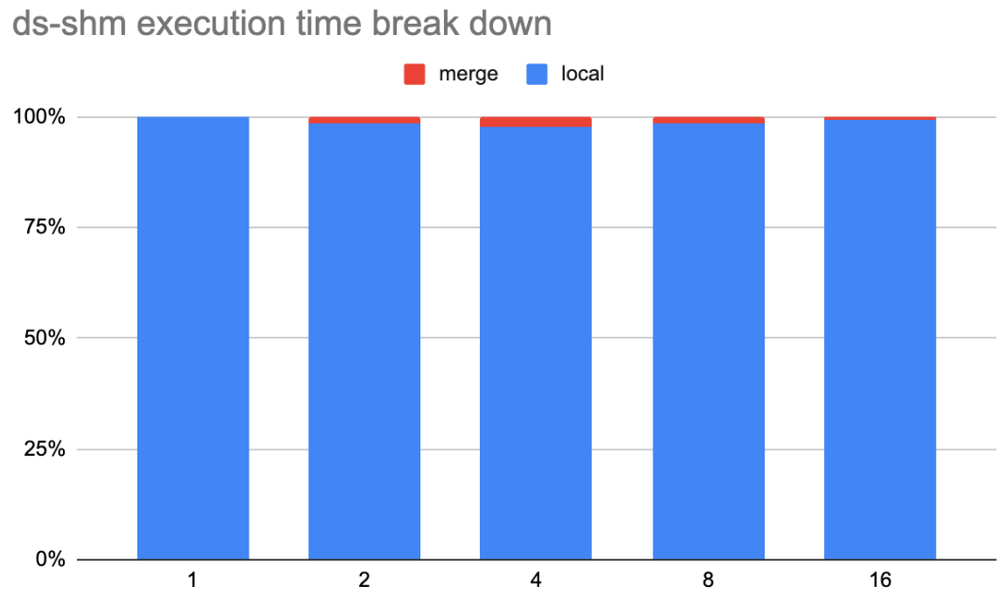

Local Computation is the bottleneck!!!

Hybrid

Key Idea: dispatch computation intensive workload to cuda,
Then leverage disjoin set for efficient merge (union)

Graph Construction (cuda) ⟶ **Disjoint Set (omp)**

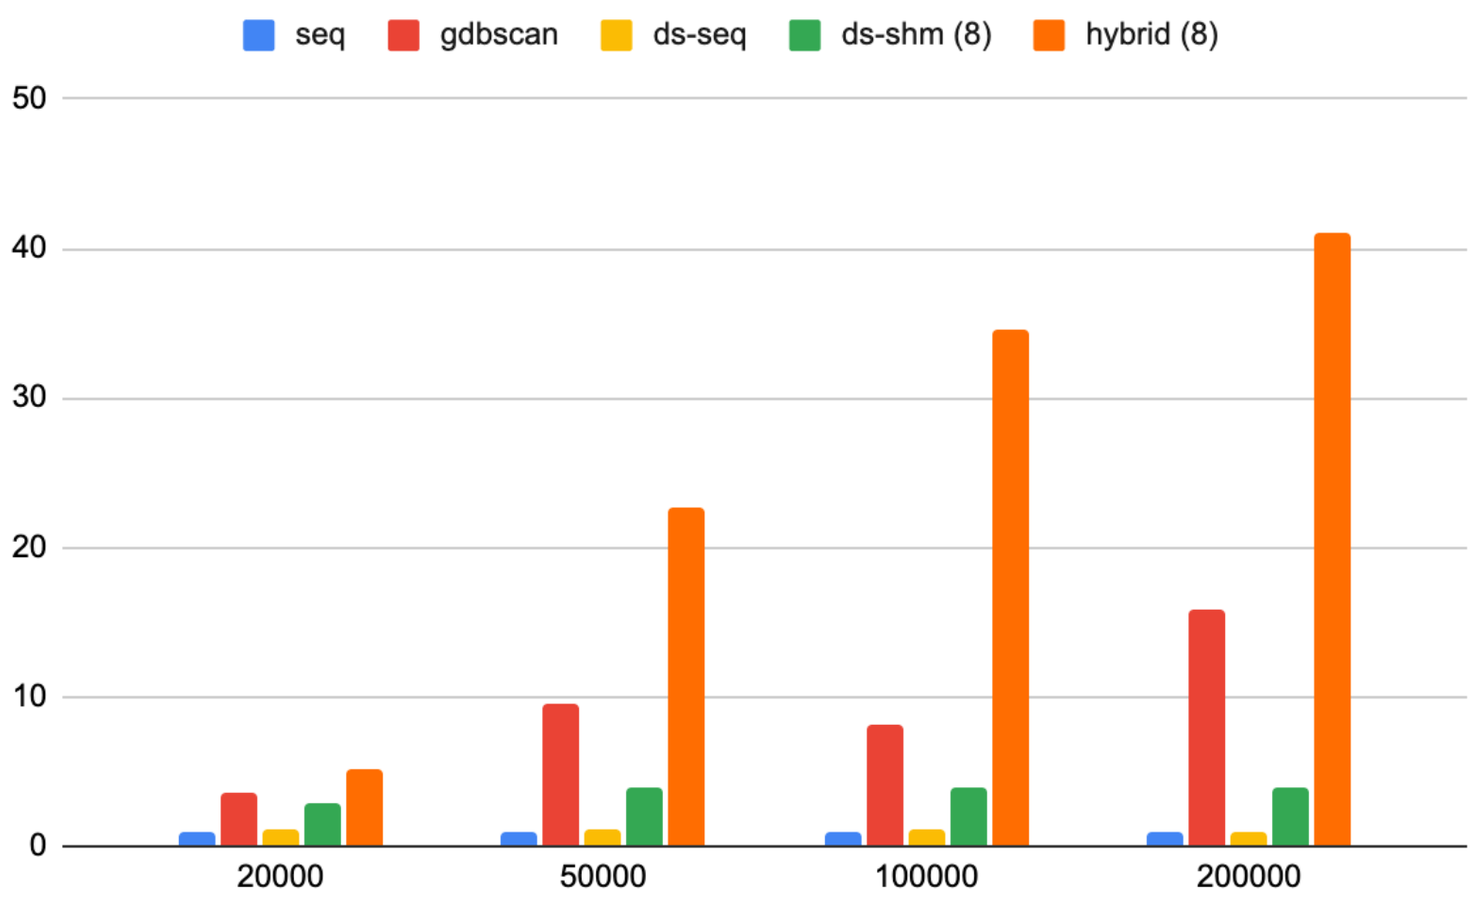

Correctness Validation